


C-space and RRTs

Textbooks




<http://planning.cs.uiuc.edu>

Configuration

- ▶ A *configuration* is a complete specification of every point of the robot
- ▶ C-space is the space of all possible configurations
- ▶ Some examples:
 - a point robot that can translate in 2D: (x, y)
 - a robot that can rotate and translate: (x, y, θ)

Example


Two Link Manipulator



Topology of Configuration Space


C-space Obstacles

Reference configuration



An obstacle in the robot's workspace


How do we get from **A** to **B**?



The C-space representation
of this obstacle...


Punching holes in the donut

- ▶ We plan for shortest paths in the C-space and not in the workspace.
- ▶ *Naïve idea*: draw a grid in the C-space
- ▶ No vertices/edges inside C-space obstacles





2D robot that can only translate

- ▶ If robot and obstacles are both polygonal, we can compute C-space obstacles by taking Minkowski difference




Non-Symmetric Robots

- ▶ If the robot is not symmetric about the origin, we should take the Minkowski difference and not the sum!
- ▶ That is “flip” the robot and then take Minkowski sum



C-space obstacles can become complicated quickly!

[Pan and Manocha, '15]




Computing exact C-space obstacles is challenging.

Solution v1

- ▶ *Idea*: checking whether a specific configuration is in collision is easier
- ▶ Let's assume we have a *fast* collision checker
 - black-box for now
- ▶ Don't construct the C-space obstacles
- ▶ Discretize the C-space
- ▶ For each vertex, check if it is in collision, else add it to the graph

Discretizing the C-space

What is a good grid resolution? Too small and we may miss solutions, too large makes it computationally expensive.



Discretizing the C-space

- ▶ Recall that we are planning in the C-space and not the workspace
- ▶ The robot is a point in C-space
- ▶ C-space can be very complex and “narrow openings”
- ▶ Finding one path, let alone, the optimal path is challenging

Solution v2

- ▶ Design a scheme that produces *sequences* of discrete samples, x_i , in the C-space

Solution v2

- ▶ Design a scheme that produces *sequences* of discrete samples, x_i , in the C-space
- ▶ Check if x_i is a collision-free configuration
- ▶ If x_i is in C_{free} then add it to the graph
 - add edges existing vertices (*more details later*)
- ▶ If x_i is not in C_{free} , then discard it
- ▶ Check if we have found a path, else repeat

Rapidly Exploring Dense Trees (RDTs)

- ▶ One of the most popular techniques
- ▶ Introduced by LaValle in '98
 - many, many, many extensions and variants

RDTs vs RRTs vs PRMs

- ▶ Many versions of the same idea
 - *with different guarantees*
- ▶ Rapidly exploring Random Trees
 - randomly sample the C-space
 - graph built will be a tree
- ▶ Rapidly exploring Dense Trees
 - any sequence of samples (not necessarily random)
- ▶ Probabilistic Roadmaps
 - build a roadmap instead of a tree
 - useful for multiple queries with diff. start and goals

And many, many variants

- ▶ articulated robots
- ▶ kinematics, dynamics, differential constraints
- ▶ <http://msl.cs.uiuc.edu/rrt/gallery.html> (circa 2000)

sampld config.

SIMPLE_RDT(q_0)

1 $\mathcal{G}.\text{init}(q_0);$

2 **for** $i = 1$ **to** k **do**

3 $\mathcal{G}.\text{add_vertex}(\alpha(i));$

4 $q_n \leftarrow \text{NEAREST}(S(\mathcal{G}), \alpha(i));$


5 $\mathcal{G}.\text{add_edge}(q_n, \alpha(i));$

starting
configura
tion


SIMPLE_RDT(q_0)

```
1   $\mathcal{G}.\text{init}(q_0);$   
2  for  $i = 1$  to  $k$  do  
3       $\mathcal{G}.\text{add\_vertex}(\alpha(i));$   
4       $q_n \leftarrow \text{NEAREST}(S(\mathcal{G}), \alpha(i));$   
5       $\mathcal{G}.\text{add\_edge}(q_n, \alpha(i));$ 
```


assume no
obstacles



*if nearest point
lies on an existing
edge, then split
that edge*



45 iterations



2345 iterations

recall that there is *no* parameter involved

RDT(q_0)

```
1   $\mathcal{G}.\text{init}(q_0);$   
2  for  $i = 1$  to  $k$  do  
3       $q_n \leftarrow \text{NEAREST}(S, \alpha(i));$   
4       $q_s \leftarrow \text{STOPPING-CONFIGURATION}(q_n, \alpha(i));$   
5      if  $q_s \neq q_n$  then  
6           $\mathcal{G}.\text{add\_vertex}(q_s);$   
7           $\mathcal{G}.\text{add\_edge}(q_n, q_s);$ 
```




Figure 5.20: If there is an obstacle, the edge travels up to the obstacle boundary, as far as allowed by the collision detection algorithm.

What about the goal?

- ▶ So far, RDT is only building a tree
- ▶ Occasionally add the goal configuration and see if it gets connected to the tree
 - say every 100th iteration

```
RDT( $q_0$ )
1   $\mathcal{G}.\text{init}(q_0);$ 
2  for  $i = 1$  to  $k$  do
3       $q_n \leftarrow \text{NEAREST}(S, \alpha(i));$ 
4       $q_s \leftarrow \text{STOPPING-CONFIGURATION}(q_n, \alpha(i));$ 
5      if  $q_s \neq q_n$  then
6           $\mathcal{G}.\text{add\_vertex}(q_s);$ 
7           $\mathcal{G}.\text{add\_edge}(q_n, q_s);$ 
```


Remember

- ▶ We are in the C-space
- ▶ A vertex in the graph is a specific configuration
- ▶ How to find the *nearest* configuration?
 - What is the *distance* function?

```
RDT( $q_0$ )
1   $\mathcal{G}.\text{init}(q_0);$ 
2  for  $i = 1$  to  $k$  do
3       $q_n \leftarrow \text{NEAREST}(S, \alpha(i));$ 
4       $q_s \leftarrow \text{STOPPING-CONFIGURATION}(q_n, \alpha(i));$ 
5      if  $q_s \neq q_n$  then
6           $\mathcal{G}.\text{add\_vertex}(q_s);$ 
7           $\mathcal{G}.\text{add\_edge}(q_n, q_s);$ 
```


Distances in C-space

- ▶ The C-space is not necessarily Euclidean space
- ▶ Need to define appropriate *distances*



Distances in C-space

- ▶ The C-space is not necessarily Euclidean space
- ▶ Need to define appropriate *distances*
- ▶ Each *edge* represents a path in C-space
 - An edge means that there is a *collision-free path* between two configurations




Practical Solutions

- ▶ Use a step size parameter
- ▶ Move in steps and check collision of a configuration
- ▶ Often, a *local steering* function is used instead*

```
BUILD_RRT( $q_{init}$ )  
1   $\mathcal{T}.\text{init}(q_{init});$   
2  for  $k = 1$  to  $K$  do  
3       $q_{rand} \leftarrow \text{RANDOM\_CONFIG}();$   
4       $\text{EXTEND}(\mathcal{T}, q_{rand});$   
5  Return  $\mathcal{T}$ 
```


```
EXTEND( $\mathcal{T}, q$ )  
1   $q_{near} \leftarrow \text{NEAREST\_NEIGHBOR}(q, \mathcal{T});$   
2  if  $\text{NEW\_CONFIG}(q, q_{near}, q_{new})$  then  
3       $\mathcal{T}.\text{add\_vertex}(q_{new});$   
4       $\mathcal{T}.\text{add\_edge}(q_{near}, q_{new});$   
5      if  $q_{new} = q$  then  
6          Return Reached;  
7      else  
8          Return Advanced;  
9  Return Trapped;
```

Figure 2: The basic RRT construction algorithm.




Tree “grows” from the start config. due to step size parameter

without step size






45 iterations



2345 iterations

with step size



▶ Is RRT
complete?

Guarantees?

Guarantees?

- ▶ Is RRT complete?
 - No but it is *probabilistically complete* – if there exists a path, the probability that RRT will not find the path decays to zero as the number of samples approaches infinity.

Guarantees?

- ▶ Is RRT complete?
 - No but it is *probabilistically complete* – if there exists a path, the probability that RRT will not find the path decays to zero as the number of samples approaches infinity.
 - Probability of failure decreases *exponentially*.
 - *Not all variants are probabilistically complete.*
- ▶ Is RRT optimal?

Guarantees?

- ▶ Is RRT optimal?
 - No. In fact, emphatically no!
 - “... *the probability that the RRT converges to an optimum solution, as the number of samples approaches infinity, is zero under some reasonable technical assumptions.*”

<https://arxiv.org/pdf/1005.0416.pdf>

Guarantees?

- ▶ Is RRT optimal?
 - No. In fact, emphatically no!
 - “... *the probability that the RRT converges to an optimum solution, as the number of samples approaches infinity, is zero under some reasonable technical assumptions.*”
- ▶ The main reason is that in RRT once we build a tree, we never modify the tree.
- ▶ RRT* is RRT + rewiring of the tree
 - Optimal!

<https://arxiv.org/pdf/1005.0416.pdf>

Why are RRTs so popular?

- ▶ Once you define the C-space and implement following subroutines, the actual algorithm is *very simple*:
 - random configuration generator
 - nearest neighbor
 - collision checker
- ▶ In *practice*, it works rather *well*.
- ▶ Can add *heuristics* on top, e.g., bias the random config. generator towards goal configuration

EXTEND() is the real crux

- ▶ EXTEND() takes the tree and *extends it closer to the given random configuration*
- ▶ We say the Euclidean case where it's a straight line
- ▶ All kinematic (and dynamic) constraints can be handled within EXTEND()

```
BUILD_RRT( $q_{init}$ )
1   $\mathcal{T}.\text{init}(q_{init});$ 
2  for  $k = 1$  to  $K$  do
3       $q_{rand} \leftarrow \text{RANDOM\_CONFIG}();$ 
4      EXTEND( $\mathcal{T}, q_{rand}$ );
5  Return  $\mathcal{T}$ 
```

```
EXTEND( $\mathcal{T}, q$ )
1   $q_{near} \leftarrow \text{NEAREST\_NEIGHBOR}(q, \mathcal{T});$ 
2  if NEW_CONFIG( $q, q_{near}, q_{new}$ ) then
3       $\mathcal{T}.\text{add\_vertex}(q_{new});$ 
4       $\mathcal{T}.\text{add\_edge}(q_{near}, q_{new});$ 
5      if  $q_{new} = q$  then
6          Return Reached;
7      else
8          Return Advanced;
9  Return Trapped;
```
