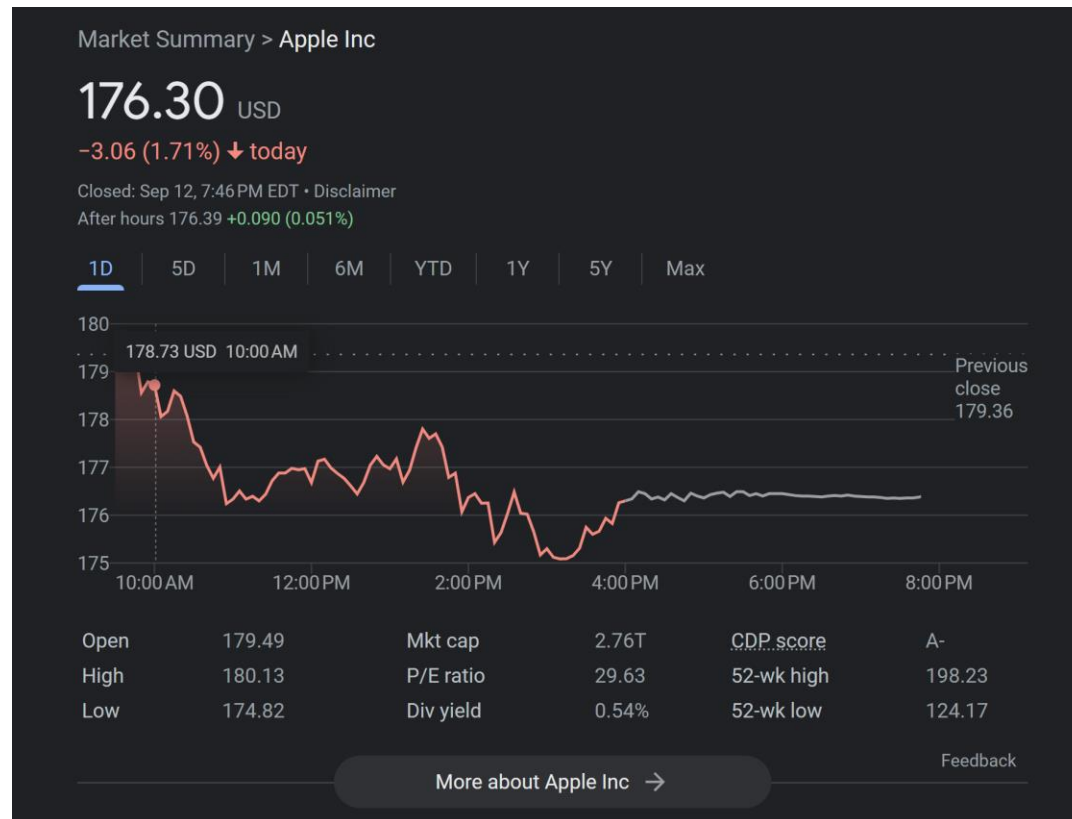

Secretary Hiring Problems & A* Algorithm

Buy one share at the lowest price

❑ Which one to buy?



The Secretary Problem

- ❑ Suppose there are n secretaries that are interviewed one at a time. The hiring decision has to be made right after an interview is done. After we interview a candidate, we know how this candidate ranks in comparison to the earlier interviewed candidates. However, we have no information about the upcoming candidates.
- ❑ *Design a strategy to hire the **best** candidate?*

The Secretary Problem

- ❑ Each secretary i has an inherent rating v_i
- ❑ But we do not know what is the maximum possible rating
- ❑ Secretaries interview in an unknown order

raise your hand if you will hire the secretary

- ❑ 1. $v_i = 10$
- ❑ 2. $v_i = 21$
- ❑ 3. $v_i = 1$
- ❑ 4. $v_i = 5$
- ❑ 5. $v_i = 15$
- ❑ 6. $v_i = 45$
- ❑ 7. $v_i = 11$
- ❑ 8. $v_i = 3$
- ❑ 9. $v_i = 2$
- ❑ 10. $v_i = 9$

raise your hand if you will hire the secretary

- ❑ 1. $v_i = 21$
- ❑ 2. $v_i = 34$
- ❑ 3. $v_i = 45$
- ❑ 4. $v_i = 5$
- ❑ 5. $v_i = 19$
- ❑ 6. $v_i = 3$
- ❑ 7. $v_i = 32$
- ❑ 8. $v_i = 4$
- ❑ 9. $v_i = 55$
- ❑ 10. $v_i = 7$

raise your hand if you will hire the secretary

- ❑ 1. $vi = 45$
- ❑ 2. $vi = 4$
- ❑ 3. $vi = 5$
- ❑ 4. $vi = 6$
- ❑ 5. $vi = 7$
- ❑ 6. $vi = 8$
- ❑ 7. $vi = 9$
- ❑ 8. $vi = 10$
- ❑ 9. $vi = 1$
- ❑ 10. $vi = 46$

The Secretary Algorithm v. 1

- ❑ Observe the first $n/2$ secretaries but don't hire anyone
- ❑ Let i be the best of the first $n/2$ secretaries
- ❑ After observing the first $n/2$ secretaries, hire the first secretary that is better than i
- ❑ If none of the remaining are better than i , then hire the last one

The Secretary Algorithm v. 1

- ❑ Worst-Case Analysis
- ❑ **Hopeless**. The worst case is when the secretaries appear in a descending order of their rank. The algorithm is arbitrarily bad. In fact, any deterministic algorithm is arbitrarily bad.

The Secretary Problem

- ❑ Suppose there are n secretaries that are interviewed one at a time. The hiring decision has to be made right after an interview is done. After we interview a candidate, we know how this candidate ranks in comparison to the earlier interviewed candidates. However, we have no information about the upcoming candidates.
- ❑ Design a strategy to **maximize** the **probability of hiring the best** candidate assuming they appear in a uniformly at random order?

Results

- If the number of secretaries tends to infinity, the **optimal** rule is to observe the first n/e secretaries, and then pick the first secretary better than the best in the first n/e secretaries

The Secretary Algorithm v. 1

- ❑ Observe the first $n/2$ secretaries but don't hire anyone
- ❑ Let i be the best of the first $n/2$ secretaries
- ❑ After observing the first $n/2$ secretaries, hire the first secretary that is better than i
- ❑ If none of the remaining are better than i , then hire the last one

- ❑ **Expected case analysis?** Can we compute the probability of picking the best candidate?

The Secretary Algorithm v. 2

- ❑ Observe the first $r-1$ secretaries but don't hire anyone
- ❑ Let i be the best of the first $r-1$ secretaries
- ❑ After observing the first $r-1$ secretaries, hire the first secretary that is better than i
- ❑ If none of the remaining are better than i , then hire the last one

The Secretary Algorithm v. 2

- ❑ Observe the first $r-1$ secretaries but don't hire anyone
- ❑ Let i be the best of the first $r-1$ secretaries
- ❑ After observing the first $r-1$ secretaries, hire the first secretary that is better than i
- ❑ If none of the remaining are better than i , then hire the last one

- ❑ **Optimal strategy:**
- ❑ Best strategy is to observe the first $\sim 37\%$ candidates and then hire the first one better than the 37%
- ❑ Probability of picking the best candidate $\sim 37\%$

Resources

- ❑ Primary: Ferguson, Thomas S. "Who solved the secretary problem?." Statistical science (1989): 282-289.

Many, many versions

- ❑ minimize the expected rank
- ❑ maximize expected rating
- ❑ ratings decay over time
- ❑ multi-choice (k choice) hiring
- ❑ submodular ratings
- ❑ matroidal constraints
- ❑ knapsack constraints
- ❑ unknown n
- ❑ sliding windows
- ❑ ...

Many, many versions

- ❑ minimize the expected rank
- ❑ maximize expected rating
- ❑ ratings decay over time
- ❑ multi-choice (k choice) hiring
- ❑ submodular ratings
- ❑ matroidal constraints
- ❑ knapsack constraints
- ❑ unknown n
- ❑ sliding windows
- ❑ ...

Maximize the expected rating

- Suppose there are n secretaries that are interviewed one at a time. The hiring decision has to be made right after an interview is done. After we interview a candidate, we know how this candidate ranks in comparison to the earlier interviewed candidates. However, we have no information about the upcoming candidates.
- Design a strategy to *maximize the expected rating* of the hired candidate assuming that the *ratings are drawn from an unknown distribution*.

Thoughts?

Maximize the expected rating

- ❑ Same strategy (Secretary Algorithm v.2)
- ❑ In the limit, the probability of picking the best candidate is $1/e$
- ❑ Therefore, the expected rating of the hired candidate is $1/e$ times the highest rating
- ❑ This is also the optimal algorithm!

Babaioff, Moshe, et al. "Online auctions and generalized secretary problems." ACM SIGecom Exchanges 7.2 (2008): 7.

Multiple Choice Secretary Problem

- Suppose there are n secretaries that are interviewed one at a time. We wish to *hire at most k* secretaries. The hiring decision has to be made right after an interview is done. Each candidate has a *rating v_i that is revealed* to us after we interview the candidate i . We have no information about the upcoming candidates.
- Design a strategy to *maximize the expected sum of ratings* of the hired candidates.

Ideas?

Build on the $k=1$ algorithm


- ❑ Observe the first $n/e - 1$ secretaries but don't hire anyone
- ❑ Let $R = \{v_1, v_2, \dots, v_k\}$ be the k best ratings observed
- ❑ After observing the first $n/e - 1$ secretaries, if a secretary has rating better than the worst rating in R , then
 - hire this secretary; and
 - remove the lowest rating in R

Build on the $k=1$ algorithm

- The previous algorithm achieves an expected rating of $1/e$ times the sum of the k largest ratings, in the limit that n tends to infinity.

Babaioff, Moshe, et al. "A knapsack secretary problem with applications." Approximation, randomization, and combinatorial optimization. Algorithms and techniques (2007): 16-28.


Revisit Dijkstra



A* Algorithm

Open and Closed Lists

- Think of open lists as frontiers of expansion
- We start with x_o and expand neighbors until we reach x_G
 - Dijkstra: expand frontier that is closest to x_o




A*

- ❑ Open list $O = \{x_o\}$, closed list $C = \{\}$
- ❑ $V(x_o) = 0$, $V(x_i) = \infty$
- ❑ repeat until x_G is in C
 - x_j : vertex in O with lowest **cost-to-come + heuristic**
 - remove it from O and store it in C , closed list
 - for each neighbor x_i of x_j not already in C
 - compute $V_{new} = C(x_j, x_i) + V(x_j)$
 - update $V(x_i)$ if it is more than V_{new}
 - if not in O , then add to O

Open and Closed Lists

- Think of open lists as frontiers of expansion
- We start with x_o and expand neighbors until we reach x_G
 - A*: expand frontier that our heuristic says will lead to minimum cost path



What's a good heuristic?

- ❑ heuristic gives an overestimate of actual $V(x_G)$
 - $h(x_j) > \text{minimum cost to reach } x_G \text{ from } x_j$
- ❑ heuristic gives an underestimate of actual $V(x_G)$
 - $h(x_j) \geq \text{minimum cost to reach } x_G \text{ from } x_j$

Admissible Heuristic

- A heuristic, $h(x_j)$, is called as admissible heuristic if and only if $h(x_j) \leq$ minimum cost to reach x_G from x_j for all x_j

- A^* will find the optimal solution as long as you have an admissible heuristic.
 - Actually, it also needs to be consistent (satisfy triangle inequality)
 - $h(x_i) \leq C(x_i, x_j) + h(x_j)$; $h(x_G) = 0$

Good Heuristics

- ❑ Needs to be an underestimate & satisfy triangle inequality to guarantee we find optimal path

- ❑ The closer your heuristic is to actual estimate, the faster you will find the optimal path
 - *need fewer expansions*

- ❑ The closer your heuristic is to 0, the slower your algorithm
 - *approaches Dijkstra; more expansions*




Admissible Heuristic

$h(x_i) = 0$ (Dijkstra)

$h(x_i)$ = Euclidean distance from x_i to x_G

...?



A*: Backtracking optimal path

$$O = \{x_O\}, C = \{\}$$

$$V(x_O) = 0, V(x_i) = \infty, \mathbf{B}(x_i)$$

repeat until x_G is in C

- x_j : vertex in O with lowest cost-to-come + heuristic
 - remove it from O and store it in C , closed list

- for each neighbor x_i of x_j not already in C
 - compute $V_{new} = C(x_j, x_i) + V(x_j)$
 - update $V(x_i)$ if it is more than V_{new} **and set $\mathbf{B}(x_i) = x_j$**
 - if not in O , then add to O

Analyzing Algorithms

1. *Completeness*

2. *Optimality*

3. *Efficiency*

Analyzing Algorithms

1. *Completeness*: Algorithm is complete if it finds the optimal solution in finite time, if a solution exists. Else, it declares failure in finite time.
2. *Optimality*: Algorithm that finds a solution whose cost is the minimum (or maximum) possible cost.
3. *Efficiency*: An algorithm is efficient if it finds the solution in the least possible time (for all inputs).

Analyzing Algorithms

1. *Completeness*: DP, Dijkstra, A^* are complete
2. *Optimality*: DP, Dijkstra, A^* are optimal.
3. *Efficiency*:
 - DP is NOT efficient.
 - Dijkstra (and A^*) is efficient if no heuristic.
 - A^* is efficient (and Dijkstra is not) with any admissible heuristic.

Optimality vs. Efficiency

- ▶ Sometimes you want a “good-enough” solution as fast as possible.
- ▶ Can we trade-off optimality and efficiency?

Weighted A*

- ▶ **Dijkstra**

- expand based on lowest $V(x_i)$

- ▶ **A***

- expand based on lowest $V(x_i) + h(x_i)$

- ▶ **ϵ Weighted A***

- expand based on lowest ???

▶ **Dijkstra**

- expand based on lowest $V(x_i)$

▶ **A***

- expand based on lowest $V(x_i) + h(x_i)$

▶ **ϵ Weighted A***


- expand based on lowest $V(x_i) + \epsilon h(x_i)$
- $\epsilon \geq 1$

Weighted A*


- $\epsilon \geq 1$
- Expand based on inflated heuristic
 - $V(x_i) + \epsilon h(x_i)$
- *Can we guarantee optimality?*

-
- ▶ $\epsilon \geq 1$
 - ▶ Expand based on inflated heuristic
 - $V(x_i) + \epsilon h(x_i)$
 - ▶ **Can we guarantee optimality?**
 - No!
 - However, we can guarantee that it will find a path whose cost is no more than ϵ times the minimum cost path.
 - $V'(x_G) \leq \epsilon V^*(x_G)$


- ▶ $\epsilon \geq 1$
- ▶ In practice, faster than A*. In fact, the higher we set ϵ to be, the faster we find the solution.



*13 expansions
solution=11 moves*



*15 expansions
solution=11 moves*



*20 expansions
solution=10 moves*

Anytime Algorithm

- ▶ If we stop the algorithm at any point in time, we should be able to return a good solution.
- ▶ More time we give an algorithm, the closer to optimal the returned solution should be.
- ▶ *Are DP, Dijkstra, A^* anytime algorithms?*

Anytime Algorithm

If we stop the algorithm at any point in time, we should be able to return a good solution.

More time we give an algorithm, the closer to optimal the returned solution should be.

Are DP, Dijkstra, A^ anytime algorithms?*

- *NO! We find the optimal path only at the last iteration.*

An anytime algorithm using weighted A^ ?*

-
- ▶ set ϵ = very high number
 - ▶ repeat until stopped
 - find path using weighted A^*
 - $\epsilon = \epsilon/2$

 - ▶ Works but we need to recompute path from scratch in every iteration. Speed up?